

KSQL / ksqlDB and Kafka Streams Interview Guide

A Light■Theme Illustrated PDF covering KSQL, KStreams, KTables, and Windowed Aggregations — with interview questions, answers, and examples.

1. ksqlDB Overview

ksqlDB is Confluent's SQL engine built on top of Kafka Streams. It enables SQL-based stream processing with STREAM and TABLE abstractions.

Concept	Description
STREAM	Unbounded, append-only event stream.
TABLE	Changelog-based view of latest state per key.
Persistent Query	Continuous transformation that runs indefinitely.
Pull Query	Snapshot query on materialized state.

```
CREATE STREAM purchases (  
  user_id VARCHAR,  
  product VARCHAR,  
  amount DOUBLE,  
  ts TIMESTAMP  
) WITH (KAFKA_TOPIC='purchases', VALUE_FORMAT='JSON', TIMESTAMP='ts');
```

```
CREATE TABLE user_spend AS  
SELECT user_id, SUM(amount) AS total  
FROM purchases  
GROUP BY user_id  
EMIT CHANGES;
```

2. Kafka Streams — KStreams and KTables

Kafka Streams is a Java API for stream processing using KStream (event stream) and KTable (state view).

Type	Description	Use Case
KStream	Continuous stream of events.	Log data, transactions.
KTable	Materialized view of latest per key.	Aggregations, changelogs.

```
StreamsBuilder builder = new StreamsBuilder();
KStream<String, String> events = builder.stream("events");
KTable<String, Long> counts = events
    .flatMapValues(v -> Arrays.asList(v.toLowerCase().split("\\W+")))
    .groupBy((k, word) -> word)
    .count(Materialized.as("word-counts"));
counts.toStream().to("output");
```

3. Stream Joins

Kafka Streams supports Stream–Stream, Stream–Table, and Table–Table joins for data enrichment and correlation.

```
-- Stream-Stream join in ksqlDB
CREATE STREAM payments (...);
CREATE STREAM shipments (...);

CREATE STREAM orders_enriched AS
SELECT p.order_id, p.amount, s.ship_status
FROM payments p
JOIN shipments s
  WITHIN 30 MINUTES
  ON p.order_id = s.order_id
EMIT CHANGES;

// Stream-Table join in Kafka Streams
KStream<String, Purchase> purchases = builder.stream("purchases");
KTable<String, User> users = builder.table("users");

KStream<String, Enriched> enriched = purchases.leftJoin(
    users,
    (p, u) -> new Enriched(p, u)
);
enriched.to("purchases-enriched");
```

4. Windowed Aggregations

Aggregations performed on events grouped by time windows (tumbling, hopping, sliding, session).

Window Type	Description	Example
Tumbling	Fixed, non-overlapping windows.	1-minute intervals.
Hopping	Overlapping windows with advance smaller than size.	5-min window hopping every 1-min.
Sliding	Window moves dynamically around each event.	10-second rolling window.
Session	Grouped by inactivity gap.	30s inactivity starts new session.

■ Tumbling Window (ksqlDB):

```
CREATE TABLE hits_per_minute AS
SELECT page_id, COUNT(*) AS hits,
       WINDOWSTART AS start_time, WINDOWEND AS end_time
FROM pageviews
WINDOW TUMBLING (SIZE 1 MINUTE)
GROUP BY page_id
EMIT CHANGES;
```

■ Hopping Window (ksqlDB):

```
CREATE TABLE avg_temp AS
SELECT sensor_id, AVG(temp) AS avg_temp
FROM readings
WINDOW HOPPING (SIZE 5 MINUTES, ADVANCE BY 1 MINUTE)
GROUP BY sensor_id
EMIT CHANGES;
```

■ Session Window (ksqlDB):

```
CREATE TABLE session_clicks AS
SELECT user_id, COUNT(*) AS clicks,
       SESSION_START() AS start_t, SESSION_END() AS end_t
FROM clicks
WINDOW SESSION (30 SECONDS)
GROUP BY user_id
EMIT CHANGES;
```

■ Tumbling Window (Kafka Streams Java):

```
KStream<String, Double> purchases = builder.stream("purchases");
purchases
    .groupByKey()
    .windowedBy(TimeWindows.of(Duration.ofMinutes(1)).grace(Duration.ofSeconds(30)))
    .reduce(Double::sum)
    .toStream()
    .to("purchases-per-minute");
```

5. Interview Notes & Best Practices

- KStream processes event streams; KTable holds the latest state per key.
- Windowed aggregations enable time-based summaries of events.
- Choose tumbling, hopping, sliding, or session windows based on analytic needs.
- Use GRACE period for late-arriving events.
- Always define event-time column (TIMESTAMP) for deterministic windows.
- All aggregations are backed by changelog topics and RocksDB state stores.
- Use EXACTLY_ONCE_V2 for transactional consistency in Kafka Streams.